

Name (Last, First): Answer Key

This practice exam covers the exercises you have done so far, 00r_hello_rust through 12c_wc. It consists of 5 questions on 21 pages; be sure you have the entire exam before starting. The point value of each question is indicated at its beginning; the entire exam has 100 points. Individual parts of a multi-part question are generally assigned approximately the same point value; exceptions are noted.

This exam is closed book. The one reference you may use is the printed course Cheatsheet. No electronic devices of any kind. You may not share material with classmates during the exam.

Be concise and clearly indicate your answer. Presentation and simplicity of your answers may affect your grade. Answer each question in the space following the question. If you find it necessary to continue an answer elsewhere, indicate clearly the location of its continuation and label its continuation with the question number and subpart if appropriate.

You should read through all the questions first, then pace yourself.

If you need to ask a question during the exam, please raise your hand and I will come to your desk. Do not come to me.

The questions begin on the next page.

Problem	Possible	Score
1	20	
2	20	
3	20	
4	20	
5	20	
Total	100	

1. (_____/20 points)

Short Questions

(a) (5 points) Exercise 01r rounds addresses up to page boundaries with these functions:

```
pub const PAGE_SIZE: usize = 0x1000;

pub fn pgroundup(addr: usize) -> usize {
    (addr + PAGE_SIZE - 1) / PAGE_SIZE * PAGE_SIZE
}

pub fn checked_pgroundup(addr: usize) -> Option<usize> {
    match addr.checked_add(PAGE_SIZE - 1) {
        Some(bumped) =>
            Some(bumped / PAGE_SIZE * PAGE_SIZE),
        None => None,
    }
}

pub fn count_pages(start: usize,
                    end: usize) -> usize {
    let mut addr = pgroundup(start);
    let mut pages = 0;
    loop {
        if addr + PAGE_SIZE > end {
            break pages;
        }
        addr += PAGE_SIZE;
        pages += 1;
    }
}
```

Give the value of each expression. Show your work.

- (i) `pgroundup(0x8000_2001)`
- (ii) `pgroundup(0x8000_5000)`
- (iii) `count_pages(0x8000_0800, 0x8000_3800)`
- (iv) `pgroundup(usize::MAX - 5)`. What happens in a debug build, and what value does it return in a release build? What does `checked_pgroundup(usize::MAX - 5)` return instead?

(i) $0x8000_2001 + 0xfff = 0x8000_3000$, a multiple of $0x1000 \rightarrow 0x8000_3000$

(ii) $0x8000_5000 + 0xfff = 0x8000_5fff$; $/$ * chops the fff $\rightarrow 0x8000_5000$
already aligned, so unchanged

(iii) `addr` starts at `pgroundup(0x8000_0800) = 0x8000_1000`
 $0x8000_1000 + 0x1000 = 0x8000_2000 \leq \text{end} \rightarrow \text{pages} = 1$
 $0x8000_2000 + 0x1000 = 0x8000_3000 \leq \text{end} \rightarrow \text{pages} = 2$
 $0x8000_3000 + 0x1000 = 0x8000_4000 > 0x8000_3800 \rightarrow \text{break} \rightarrow 2$
 (the ragged $0x800$ bytes at each end are not whole pages)

(iv) debug: `addr + 0xfff` overflows $\rightarrow$ panic: "attempt to add with overflow"

release: the add wraps. $2^{64} - 6 + 4095 = 2^{64} + 4089 \rightarrow 4089$,

and $4089 / 4096 * 4096 = 0$ — "round up" went down to 0!

`checked_pgroundup`: `checked_add` fails $\rightarrow$ None

(b) (5 points) In exercise 04r, Pte derives Copy and PageGuard implements Drop:

```
#[derive(Clone, Copy, PartialEq, Eq, Debug)]
pub struct Pte(pub u64);          // pa(self) and flags(self) take self BY VALUE

pub struct PageGuard<'a> {
    free: &'a mut Vec<usize>,
    pa: usize,
}
// PageGuard::take(free: &'a mut Vec<usize>) -> PageGuard<'a>
// pa(&self) -> usize    entry(&self, flags: u64) -> Pte    release(self) {}
// PTE_V and PTE_R are u64 flag bits, 1 and 2.

impl Drop for PageGuard<'_> {
    fn drop(&mut self) {
        self.free.push(self.pa);
    }
}
```

Consider this code:

```
let mut free = vec![0x8000_0000, 0x8000_1000];
let page = PageGuard::take(&mut free);    // 1
let pte = page.entry(PTE_V | PTE_R);      // 2
let a = pte.pa();                         // 3
let b = pte.flags();                      // 4
let left = free.len();                    // 5
page.release();                           // 6
let c = page.pa();                        // 7
let d = free.len();                       // 8
```

Two of the numbered lines do not compile. Which two, and why? Why does line 4 compile, when line 3 already passed pte to a method that takes self by value? Why can PageGuard never derive Copy?

Line 5 does not compile: page holds the exclusive &mut free borrow, and page is still used on line 6, so free can't be read while it is lent out. (E0502)

Line 7 does not compile: release(self) took page by value, so page was moved into release and dropped at its closing brace (the page went back on the list). Using page afterwards is a use of a moved value. (E0382)

Line 4 is fine because Pte is Copy: line 3 copied the 8 bytes into pa(), so pte is still ours. (Line 8 is fine too: the borrow of free ended when page died inside release.)

Copy + Drop: every copy of a guard would run drop, so the same page would be pushed back twice = a double free. rustc refuses (E0184).

(c) (5 points) Exercise 03r's `longest` returns the longer of two slices. Does the code below compile? If not, explain what the compiler objects to, given that at run time `r` would point at `a`, the longer slice, which is still alive. Show one way to fix it.

```
pub fn longest<'a>(a: &'a [u8], b: &'a [u8]) -> &'a [u8] {  
    if a.len() >= b.len() { a } else { b }  
}  
  
let a = [1u8, 2, 3];  
let r;  
{  
    let b = [4u8, 5];  
    r = longest(&a, &b);  
}  
println!("{}", r.len());
```

No. E0597: 'b' does not live long enough.

The signature gives the result the same lifetime 'a as both inputs, so as far as the compiler knows `r` may borrow from `b`. It checks the signature, not which branch runs. `b` dies at the `}` but `r` is used after it $\rightarrow$ rejected.

Fix: move the `println!` inside the block, or declare `b` in the outer scope next to `a`.

(d) (5 points) Exercise 08r builds a small filesystem. Its directory holds three entries: `README` and `init` are files (`init` contains the 18 bytes `first user process`), and `dev` is a directory. This is its code; the questions are on the next page.

```
pub const NAME_MAX: usize = 14;

#[derive(Debug, Clone, Copy, PartialEq, Eq)]
pub enum FsError { NotFound, NameTooLong, IsADirectory }

#[derive(Debug, Clone, Copy, PartialEq, Eq)]
pub enum InodeKind { File, Dir }

#[derive(Debug, Clone, Copy, PartialEq, Eq)]
pub struct Inode {
    pub inum: u32,           // the index of this inode's bytes in `data`
    pub kind: InodeKind,
    pub size: usize,
}

pub struct DirEnt {
    pub name: String,
    pub inode: Inode,
}

pub struct FileSystem {
    entries: Vec<DirEnt>,    // the directory: one entry per name
    data: Vec<Vec<u8>>,      // data[i] is the contents of the inode whose inum is i
}

impl FileSystem {
    pub fn find(&self, name: &str) -> Option<Inode> {
        for entry in self.entries.iter() {
            if entry.name == name {
                return Some(entry.inode);
            }
        }
        None
    }

    pub fn lookup(&self, name: &str) -> Result<Inode, FsError> {
        if name.len() > NAME_MAX {
            return Err(FsError::NameTooLong);
        }
        self.find(name).ok_or(FsError::NotFound)
    }

    pub fn read(&self, inode: Inode) -> Result<&[u8], FsError> {
        if inode.kind == InodeKind::Dir {
            return Err(FsError::IsADirectory);
        }
        Ok(&self.data[inode.inum as usize])
    }

    pub fn read_file(&self, name: &str) -> Result<&[u8], FsError> {
        let inode = self.lookup(name)?;
        self.read(inode)
    }
}
```

(d), continued. The filesystem code is on the previous page.

- (i) Rewrite `read_file` without the `?` operator.
- (ii) For each call below, give the value `read_file` returns and say which function produced it. The last name is 16 characters long.

`read_file("init")` `read_file("dev")` `read_file("sixteen_chars_xx")`

(i)

```
pub fn read_file(&self, name: &str) -> Result<&[u8], FsError> {
    let inode = match self.lookup(name) {
        Ok(inode) => inode,
        Err(e) => return Err(e),    // what ? does
    };
    self.read(inode)
}
```

(ii) `"init"` → `Ok(b"first user process")`, from `read` (lookup was `Ok`)

`"dev"` → `Err(IsADirectory)`, from `read` (lookup found `dev`)

`"sixteen_chars_xx"` → `Err(NameTooLong)`, from `lookup` (`16 > 14`); the early return passes it straight up and `read` never runs

2. (_____/20 points)

Process States

Exercise 05r models a process's life as an enum. This is its solution, with the comments shortened. The blank lines are for part (b).

```
#[derive(Debug, Clone, Copy, PartialEq, Eq)]
pub enum ProcState {
    Unused,
    Runnable,
    Running,
    Sleeping { channel: u64 },
    Zombie { exit_status: i32 },
}

#[derive(Debug, Clone, Copy, PartialEq, Eq)]
pub enum Event {
    Alloc, // the kernel hands out this table slot
    Schedule, // the scheduler picks this process
    Preempt, // the timer takes the CPU back
    Sleep { channel: u64 }, // block until someone wakes this channel
    Wake { channel: u64 }, // a device or another process signals a channel
    Exit(i32), // the process finishes, with a status
    Reap, // the parent collects the corpse, freeing the slot
    Kill,
}

impl ProcState {
    pub fn can_run(&self) -> bool {
        match self {
            ProcState::Runnable | ProcState::Running => true,
            ProcState::Unused => false,
            ProcState::Sleeping { .. } => false,
            ProcState::Zombie { .. } => false,
        }
    }
}

pub fn next_state(state: ProcState, event: Event) -> Option<ProcState> {
    match (state, event) {
        (ProcState::Unused, Event::Alloc) => Some(ProcState::Runnable),
        (ProcState::Runnable, Event::Schedule) => Some(ProcState::Running),
        (ProcState::Running, Event::Preempt) => Some(ProcState::Runnable),
        (ProcState::Running, Event::Sleep { channel }) => Some(ProcState::Sleeping { channel }),
        (ProcState::Sleeping { channel: c }, Event::Wake { channel: w }) if c == w => {
            Some(ProcState::Runnable)
        }
        (ProcState::Running, Event::Exit(exit_status)) => Some(ProcState::Zombie { exit_status }),
        (ProcState::Zombie { .. }, Event::Reap) => Some(ProcState::Unused),
        (ProcState::Runnable | ProcState::Running | ProcState::Sleeping { .. }, Event::Kill)
            => Some(ProcState::Zombie { exit_status: -1 }),
        // must go above the _ arm
        _ => None,
    }
}

pub fn apply(state: ProcState, event: Event) -> ProcState {
    match next_state(state, event) {
        Some(new_state) => new_state,
        None => state,
    }
}
```

(a) (6 points) Starting from `ProcState::Unused`, the kernel runs `state = apply(state, e)` for each event `e` below, in order. Give the state after each call, and mark each event that `apply` ignored because `next_state` returned `None`.

Step	Event <code>e</code>	State after the call	Ignored?
1	Alloc	<i>Runnable</i>	
2	Schedule	<i>Running</i>	
3	Sleep { channel: 0x10 }	<i>Sleeping { channel: 0x10 }</i>	
4	Wake { channel: 0x20 }	<i>Sleeping { channel: 0x10 }</i>	<i>yes, wrong channel</i>
5	Schedule	<i>Sleeping { channel: 0x10 }</i>	<i>yes, not Runnable</i>
6	Wake { channel: 0x10 }	<i>Runnable</i>	
7	Preempt	<i>Runnable</i>	<i>yes, not Running</i>
8	Schedule	<i>Running</i>	
9	Exit(3)	<i><u>Zombie { exit_status: 3 }</u></i>	

(b) (7 points) Add a new event, `Event::Kill`. Killing a process that is `Runnable`, `Running`, or `Sleeping` makes it `Zombie { exit_status: -1 }`. Killing an `Unused` slot or a `Zombie` is illegal. Write the new variant and *one* new match arm in the blank lines of the listing on the previous page. Then explain why it matters where your arm goes relative to `_ => None`, and why its position relative to the other arms does not.

(Kill variant and the arm are written in the listing.)

Arms are tried top to bottom and the first match wins. `_` matches every pair, so an arm below it could never run: Kill would silently stay illegal (rustc only warns “unreachable pattern”).

No other arm mentions `Event::Kill`, so none of them overlaps the new one; above or below them, it gives the same answer.

(c) (7 points) Now add a state `ProcState::Stopped` (a process paused by the user, as with Ctrl-Z) and two events: `Stop` moves a `Running` process to `Stopped`, and `Continue` moves a `Stopped` process to `Runnable`.

- (i) You add just the three new variants to the two enums and recompile. Which functions in the listing stop compiling, and which still compile? Explain why, in terms of exhaustiveness. Why is the function that still compiles the more dangerous one?
- (ii) Write the code you must add to finish the change.

(i) can_run stops compiling: E0004, non-exhaustive patterns: `&ProcState::Stopped` not covered. It has one arm per variant, so the new variant has no arm.

next_state still compiles (and so do `apply` and the enums): its `_ => None` arm covers every pair, including all the new ones.

That is the danger: the compiler cannot remind you that `Stop` and `Continue` need arms. Until you remember, both are silently illegal. `can_run`'s one-arm-per-variant style is what makes the compiler hand you the list.

(ii)

```
// in can_run:
ProcState::Stopped => false,
```

```
// in next_state, above the _ arm:
(ProcState::Running, Event::Stop) => Some(ProcState::Stopped),
(ProcState::Stopped, Event::Continue) => Some(ProcState::Runnable),
```

(Optional: add `ProcState::Stopped` to the `Kill` arm from (b), so a stopped process can be killed too.)

3. (_____/20 points)

Scheduling Policies as Traits

Exercise 07r separates the scheduling *policy* from the loop that uses it:

```
#[derive(Clone, Copy, PartialEq, Eq, Debug)]
pub enum ProcState { Unused, Runnable, Sleeping }

pub trait Scheduler {
    fn pick_next(&mut self, states: &[ProcState]) -> Option<usize>;

    fn run_for(&mut self, states: &[ProcState], turns: usize) -> Vec<usize> {
        let mut picked = Vec::new();
        for _ in 0..turns {
            match self.pick_next(states) {
                Some(slot) => picked.push(slot),
                None => break,
            }
        }
        picked
    }
}

pub struct RoundRobin {
    pub next: usize,
}

impl RoundRobin {
    pub fn new() -> RoundRobin {
        RoundRobin { next: 0 }
    }
}

impl Scheduler for RoundRobin {
    fn pick_next(&mut self, states: &[ProcState]) -> Option<usize> {
        let n = states.len();
        for off in 0..n {
            let slot = (self.next + off) % n;
            if states[slot] == ProcState::Runnable {
                self.next = (slot + 1) % n;
                return Some(slot);
            }
        }
        None
    }
}
```

(a) (6 points) What does `picked` hold after the code below, and what is `rr.next`? Show which slots each call to `pick_next` looks at.

```
use ProcState::*;  
let states = [Sleeping, Runnable, Runnable, Unused, Runnable];  
let mut rr = RoundRobin::new();  
let picked = rr.run_for(&states, 5);
```

call 1: `next = 0`: looks at 0 (Sleeping), 1 (Runnable) $\rightarrow$ 1, `next = 2`

call 2: looks at 2 $\rightarrow$ 2, `next = 3`

call 3: looks at 3 (Unused), 4 $\rightarrow$ 4, `next = (4+1) % 5 = 0`

call 4: looks at 0, 1 $\rightarrow$ 1, `next = 2`

call 5: looks at 2 $\rightarrow$ 2, `next = 3`

`picked` = [1, 2, 4, 1, 2] `rr.next` = 3

(b) (6 points) The 07r file also has a second trait, `Out`, which two types implement, a second `Scheduler`, and two functions that use them:

```
pub trait Out {
    fn write_str(&mut self, s: &str);
    fn write_line(&mut self, s: &str) {
        self.write_str(s);
        self.write_str("\n");
    }
}

impl Out for StringOut { ... } // keeps everything written in a String
impl Out for CountingOut { ... } // only counts the bytes and the calls
impl Scheduler for Priority { ... } // runs the most urgent runnable slot

pub fn write_banner(out: &mut dyn Out) {
    out.write_line("rv6 kernel ready");
}

pub fn trace<S: Scheduler, O: Out>(sched: &mut S, states: &[ProcState],
                                   names: &[&str], turns: usize, out: &mut O) {
    for slot in sched.run_for(states, turns) {
        out.write_line(names[slot]);
    }
}
```

A program calls `trace` four times, with the type pairs $(S, O) = (\text{RoundRobin}, \text{StringOut})$, $(\text{Priority}, \text{StringOut})$, $(\text{RoundRobin}, \text{StringOut})$, and $(\text{Priority}, \text{CountingOut})$. It calls `write_banner` twice, once with a `StringOut` and once with a `CountingOut`.

- (i) How many copies of `trace`, and how many of `write_banner`, does the compiler put in the binary? Name the kind of dispatch each one uses and say how it finds the right method.
- (ii) What goes wrong if you delete `: Scheduler` from `trace`'s type parameters?

(i) `trace`: 3 copies, one per distinct (S, O) pair; the repeated $(\text{RoundRobin}, \text{StringOut})$ reuses its copy. This is monomorphization, i.e. static dispatch: each call is resolved at compile time and can be inlined.

`write_banner`: 1 copy, dynamic dispatch. The `&mut dyn Out` carries a vtable pointer, and the right `write_str` is looked up at run time.

(ii) E0599: no method named `run_for` found for `&mut S`. Without the bound the body may only assume `S` is some type, and the bound is the only thing that promises `S` has `run_for`.

(c) (8 points) Write `Backward`, a policy that rotates the other way. Each call scans *downward* from `self.next` (`next`, `next - 1`, ..., `0`, then wrapping around to the last slot), takes the first `Runnable` slot, and leaves the cursor one slot below the one it took (wrapping from `0` to the last slot). Like `RoundRobin`, it must return `None` when nothing is runnable, including for an empty table, and it must never panic. Then give what `Backward::new().run_for(&states, 5)` returns for the table in part (a).

```
pub struct Backward {
    pub next: usize,
}

impl Backward {
    pub fn new() -> Backward {
        Backward { next: 0 }
    }
}

impl Scheduler for Backward {
    fn pick_next(&mut self, states: &[ProcState]) -> Option<usize> {

        let n = states.len();
        for off in 0..n {
            // + n first: self.next - off would underflow
            // and panic in a debug build (O1r) once off > next
            let slot = (self.next + n - off) % n;
            if states[slot] == ProcState::Runnable {
                self.next = (slot + n - 1) % n;
                return Some(slot);
            }
        }
        None // also covers n == 0: the loop never runs

        run_for(&states, 5) = [4, 2, 1, 4, 2]
        scans: 0, 4 → 4; 3, 2 → 2; 1 → 1; 0, 4 → 4; 3, 2 → 2

    }
}
```

4. (_____/20 points)

Streaming with O(1) State

This is `count_fd` from the 12c wc solution. `ulib::read(fd, &mut buf)` reads up to `buf.len()` bytes into `buf` and returns `Ok(n)`, the number it read, which may be fewer than asked for; `Ok(0)` means end of file.

```
#[derive(Default, Clone, Copy)]
struct Counts { lines: usize, words: usize, bytes: usize }

fn is_space(b: u8) -> bool {
    matches!(b, b' ' | b'\t' | b'\n' | b'\r' | 0x0b | 0x0c)
}

fn count_fd(fd: Fd) -> Result<Counts, ulib::Error> {
    let mut c = Counts::default();
    let mut in_word = false;
    let mut buf = [0u8; 512];
    loop {
        let n = ulib::read(fd, &mut buf)?;
        if n == 0 {
            return Ok(c);
        }
        for &b in &buf[..n] {
            c.bytes += 1;
            if b == b'\n' {
                c.lines += 1;
            }
            if is_space(b) {
                in_word = false;
            } else if !in_word {
                in_word = true;
                c.words += 1;
            }
        }
    }
}
```

(a) (6 points) `count_fd(STDIN)` runs, and its calls to `ulib::read` return the bytes below (`␣` is a space, `\n` a newline). Fill in the table with the values after each call has been handled, then give what `count_fd` returns.

Call	read returns	n	c.lines	c.words	c.bytes	in_word
1	ab␣c	4	0	2	4	true
2	d\n␣␣e	5	1	3	9	true
3	f\n	2	2	3	11	false
4	end of file	0	2	3	11	false

`count_fd` returns:

`Ok(Counts { lines: 2, words: 3, bytes: 11 })`

(b) (4 points) A student moves `let mut in_word = false;` from above the loop to the first line *inside* it. All five of `wc`'s tests still pass. Their inputs are `one two\nthree\n`, `aaabbb\n`, no newline **here**, the empty input, and two files holding `x\n` and `y z\n`. The test harness answers each `read` with as many of the remaining bytes as fit in the buffer. What does the student's `count_fd` return for the input in part (a)? Why do the tests not catch it?

lines 2, words 5, bytes 11. `in_word` is forgotten at every read, so the `d` at the start of call 2 (the end of "`cd`") and the `f` at the start of call 3 (the end of "`ef`") each count as a new word.

Every test input is under 512 bytes, so the harness returns it whole in one read and no word is ever split across reads. Only input that arrives in several reads (a file over 512 bytes, a pipe, a terminal) shows the bug.

(c) (10 points) Using `count_fd` as your model, write

```
fn longest_line(fd: Fd) -> Result<usize, ulib::Error>
```

which returns the length in bytes of the longest line read from `fd`, not counting its `\n`. A last line with no trailing newline still counts, and empty input gives 0. For example, `hello\nworld!!\n` gives 7 and `a\n\nbcd` gives 3. Follow `wc`'s rules: one fixed 512-byte buffer, no `Vec` or `String`, and a fixed number of variables however large the input is. Lines can be longer than 512 bytes, and a line can be split across reads.

```
fn longest_line(fd: Fd) -> Result<usize, ulib::Error> {
    let mut buf = [0u8; 512];
    let mut cur = 0;    // the line we are in, so far
    let mut longest = 0; // the longest line seen
    loop {
        let n = ulib::read(fd, &mut buf)?;
        if n == 0 {
            if cur > longest {    // a last line with no '\n'
                longest = cur;
            }
            return Ok(longest);
        }
        for &b in &buf[..n] {
            if b == b'\n' {
                if cur > longest {
                    longest = cur;
                }
                cur = 0;
            } else {
                cur += 1;
            }
        }
    }
}
```

`cur` and `longest` live outside the loop, like `in_word`, so a line split across reads keeps counting. Two `usizes`, whatever the input size: $O(1)$.

5. (_____/20 points)

The Commands

The command exercises use these ulib calls and Args methods:

```
args.len()           // number of arguments, including argv[0], the command's name
args.get(i)          // Option<&[u8]>: argument i, or None if i >= args.len()
ulib::read(fd, &mut buf) // Result<usize, Error>: reads up to buf.len() bytes; Ok(0) = end of file
ulib::write_all(fd, bytes) // Result<(), Error>: writes every byte of bytes
ulib::open(path, O_RDONLY) // Result<Fd, Error>: Err if the file cannot be opened
ulib::close(fd)          // gives the file descriptor back
ulib::write_usize(fd, n, w) // writes n in decimal, right-aligned in w columns
```

(a) (8 points) Many versions of `echo` accept the option `-n`: when the *first* argument is exactly `-n`, `echo` does not print it and does not print the trailing newline. So `echo -n a b` prints `a b` with no newline, `echo -n` prints nothing at all, and `echo a -n` prints `a -n` and a newline. Here is the 10c solution's run. Modify it to support `-n`, using the blank lines, and clearly mark any line you change.

```
fn run(args: Args) -> i32 {
    let mut first = 1;
    let mut newline = true;
    if args.get(1) == Some(&b"-n"[..]) {
        first = 2;
        newline = false;
    }

    for i in 1'first..args.len() {
        if i > 1'first {           ← else echo -n a b prints a space before a
            let _ = ulib::write_all(STDOUT, b" ");
        }
        let _ = ulib::write_all(STDOUT, args.get(i).unwrap());
    }
    if newline {
        let _ = ulib::write_all(STDOUT, b"\n");
    }
    0
}
```

The separator rule has to follow the new first argument: a space goes between printed arguments, so the test is `i > first`, not `i > 1`. (Also accepted: `args.len() > 1 && args.get(1).unwrap() == b"-n".`)

(b) (6 points) The 11c solution's `cat_fd` is on the left. A student writes the version on the right, with no loop.

```
fn cat_fd(fd: Fd) -> Result<(), ulib::Error> {
    let mut buf = [0u8; 512];
    loop {
        let n = ulib::read(fd, &mut buf)?;
        if n == 0 {
            return Ok(()); // end of file
        }
        ulib::write_all(STDOUT, &buf[..n])?;
    }
}

fn cat_fd(fd: Fd) -> Result<(), ulib::Error> {
    let mut buf = [0u8; 512];
    let n = ulib::read(fd, &mut buf)?;
    ulib::write_all(STDOUT, &buf[..n])?;
    Ok(())
}
```

- (i) `cat`'s six tests check that it copies a file, concatenates several files, reads standard input when given no arguments, reports a missing file, keeps going after a bad file, and copies a 1500-byte file. Every other test file is under 100 bytes, and the test harness answers each `read` with as many of the remaining bytes as fit in the buffer. The student's version passes five of the six tests. Which one fails, and why?
- (ii) With the correct `cat_fd`, how many calls to `ulib::read` and how many to `ulib::write_all` does it take to copy a 1300-byte file, if every `read` returns as many bytes as fit in the buffer?
- (iii) The same file, but now the operating system's `read` never returns more than 128 bytes per call. How many calls of each? Why is a short read like this not an error?

(i) the 1500-byte file: one read returns at most 512 bytes, so only the first 512 are copied. The small files fit in one read, so they pass.

(ii) reads return 512, 512, 276, 0 $\rightarrow$ 4 reads, 3 write_all

(iii) $1300 = 10 \times 128 + 20 \rightarrow$ 11 reads with data + the 0 = 12 reads,

11 write_all. read returns how many bytes it did read; fewer than asked is normal (a file's end, a terminal line, a pipe). Only 0 means end of file.

(c) (6 points) This is the rest of the 12c wc solution. Counts, is_space, and count_fd are in Problem 4. The question is on the next page.

```
fn report(c: Counts, label: Option<&[u8]>) {
    let _ = ulib::write_usize(STDOUT, c.lines, 8);
    let _ = ulib::write_usize(STDOUT, c.words, 8);
    let _ = ulib::write_usize(STDOUT, c.bytes, 8);
    if let Some(name) = label {
        let _ = ulib::write_all(STDOUT, b" ");
        let _ = ulib::write_all(STDOUT, name);
    }
    let _ = ulib::write_all(STDOUT, b"\n");
}

fn run(args: Args) -> i32 {
    if args.len() < 2 {
        return match count_fd(STDIN) {
            Ok(c) => {
                report(c, None);
                0
            }
            Err(_) => 1,
        };
    }

    let mut status = 0;
    let mut total = Counts::default();
    let multiple = args.len() > 2;

    for i in 1..args.len() {
        let path = args.get(i).unwrap();
        match ulib::open(path, O_RDONLY) {
            Ok(fd) => {
                match count_fd(fd) {
                    Ok(c) => {
                        report(c, Some(path));
                        total.lines += c.lines;
                        total.words += c.words;
                        total.bytes += c.bytes;
                    }
                    Err(_) => status = 1,
                }
                let _ = ulib::close(fd);
            }
            Err(_) => {
                let _ = ulib::write_all(STDERR, b"wc: cannot open ");
                let _ = ulib::write_all(STDERR, path);
                let _ = ulib::write_all(STDERR, b"\n");
                status = 1;
            }
        }
    }
    if multiple {
        report(total, Some(b"total"));
    }
    status
}
```

(c), continued. `a.txt` holds the 22 bytes `one\two\tthree\n\tfour\n`. `b.txt` holds the 21 bytes `last line, no newline`, with no newline at the end. `nosuch` does not exist. Using the code on the previous page, fill in what this command prints:

```
$ wc a.txt nosuch b.txt
```

	lines	words	bytes	label
stdout, line 1	2	4	22	a.txt
stdout, line 2	0	4	21	b.txt
stdout, line 3	2	8	43	total
stdout, line 4	(none)			

Standard error: `wc: cannot open nosuch`

Exit status (\$?): 1

Not every line of the table need be used. Then say how the output changes for `wc b.txt`, and for `wc < b.txt`.

The total row is printed because three files were named, even though only two opened; the failure goes to stderr and makes the status 1.

`wc b.txt`: just `0 4 21 b.txt`, with no total row (one file), exit 0.

`wc < b.txt`: `0 4 21` with no label, since `wc` read standard input and has no name for it; exit 0.

Continue your answers here if necessary.